

## Training Phase

**Objective Function:** The model minimizes a loss function, often the **cross-entropy loss**:

$$L(\theta) = - \sum_{i=1}^N y_i \log(p_{\theta}(y_i|x_i))$$

Where:

- $\theta$  represents the model parameters.
- $x_i$  is the input sequence.
- $y_i$  is the target (true next token).
- $p_{\theta}(y_i|x_i)$  is the probability the model assigns to  $y_i$  given  $x_i$ .

**Gradient Descent:** Parameters are updated using gradient descent:

$$\theta_{t+1} = \theta_t - \alpha \nabla L(\theta_t)$$

- $\alpha$  is the learning rate.
- $\nabla L(\theta_t)$  is the gradient of the loss with respect to the parameters.

**Backpropagation:** The computation of gradients involves backpropagation through the neural network, updating parameters layer by layer.

## Inference Phase

**Prediction:** For a given input sequence  $x$ , the model predicts the next token:

$$\hat{y} = \operatorname{argmax}_y p_{\theta}(y|x)$$

Where  $\hat{y}$  is the predicted next token.

**Sequence Generation:** For generating longer sequences, this process is repeated:

$$y_1, y_2, \dots, y_T = \operatorname{argmax}_{y_1, \dots, y_T} \prod_{t=1}^T p_{\theta}(y_t|x, y_1, \dots, y_{t-1})$$

Here, each token's prediction depends on the previous tokens and the input.

## Few-Shot and Zero-Shot Learning

**Few-Shot:** The model leverages conditional probability for task-specific examples:

$$P(Y|X, E) = \prod_i P(y_i|x, e_1, ..., e_n, y_1, ..., y_{i-1})$$

Where  $E = \{e_1, ..., e_n\}$  are example pairs.

**Zero-Shot:** The model infers from its training distribution without specific examples:

$$P(Y|X) \approx \sum_{E \in \text{training set}} P(Y|X, E)$$

## Fine-Tuning

**Parameter Adjustment:** Fine-tuning involves adjusting parameters with a new loss function tailored to specific tasks:

$$L_{\text{fine-tune}}(\theta) = - \sum_{i=1}^M \tilde{y}_i \log(p_{\theta}(\tilde{y}_i|\tilde{x}_i))$$

Where:

- $\tilde{x}_i, \tilde{y}_i$  are from a task-specific dataset.
- $M$  is the number of examples in the fine-tuning set.

This mathematical framework helps understand how LLMs are initially trained, how they generate responses, and how they can be adapted or fine-tuned for specific applications without real-time learning during user interactions.